



Multi-Process Programming in C

Advanced Operating Systems (2016/2017)

Giuseppe Massari
giuseppe.massari@polimi.it

Multi-process programming

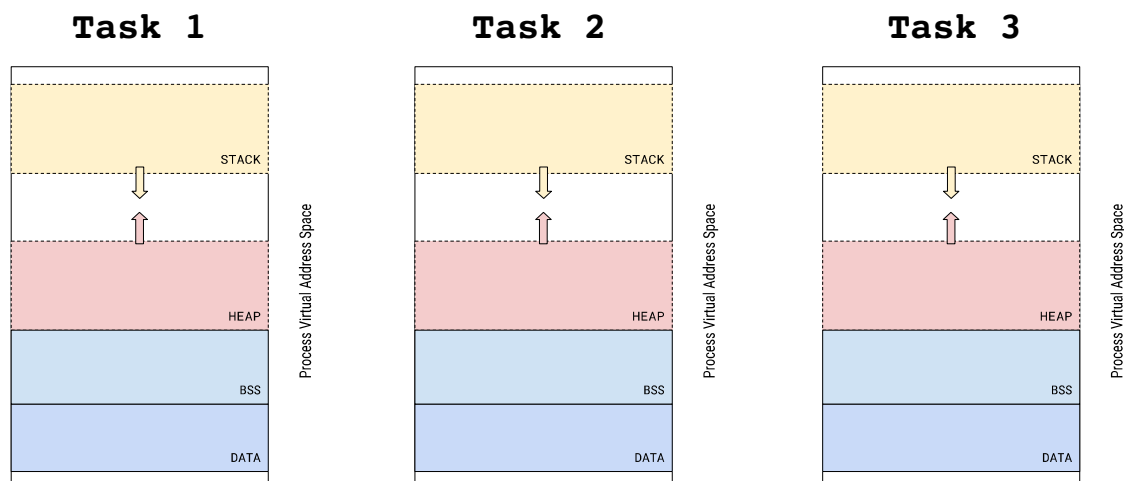
- Fork processes
- Inter-process synchronization
- Executing other programs

Inter-Process Communication

- Signals
- Pipes
- Shared memory
- Synchronization

Why multi-process programming?

- Multi-process means that each task has its own *address space*
 - ♦ More *task isolation* and independence compared to multi-threading
- Useful choice for multi-tasking application where tasks have significant requirements in terms of *resources*
 - ♦ Tasks requiring “long” processing times
 - ♦ Tasks processing big data structures
 - ♦ Tasks featuring high I/O activity (networking, disk accesses, ...)



Example 1: *Forking a process*

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main () {
    pid_t child_pid;
    printf("Main process id = %d (parent PID = %d)\n",
        (int) getpid(), (int) getppid());

    child_pid = fork();
    if (child_pid != 0)
        printf("Parent: child's process id = %d\n", child_pid);
    else
        printf("Child: my process id = %d\n", (int) getpid());
    return 0;
}
```

- **fork()** creates a new process duplicating the calling process

Example 1: *Forking a process*

```
$ gcc example1.c -o fork_ex1  
$ ./fork_ex1
```

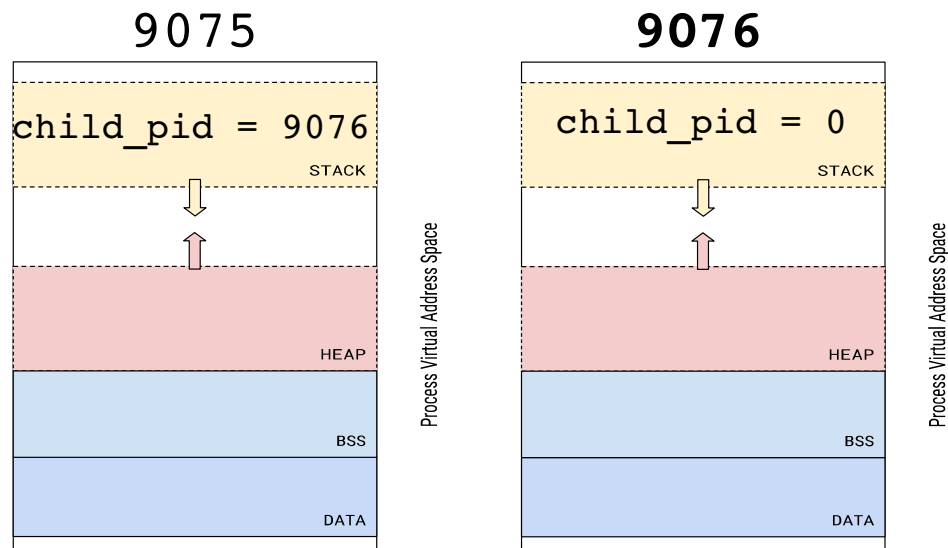
```
Main process id = 9075 (parent PID = 32146)  
Parent: child's process id = 9076  
Child:  my process id = 9076
```

- The main process has PID = 9075. It's parent (PID=32146) is the shell (echo \$\$) from which the executable has been started
- After the **fork()** the program concurrently executes two processes
- The `child_pid` variable, in the address space of the *parent process*, is set to the return value of the fork (the child process ID)
- The `child_pid` variable, in the address space of the *child process*, is not set
- The **getpid()** returns the current process identifying number

Example 1: *Forking a process*

```
int main ()
{
    pid_t child_pid;
    ...

    child_pid = fork();
    ...
}
```



- Parent process virtual address space is replicated in the child
 - ♦ Including the states of variables, mutexes, condition variables, POSIX objects
- The child inherits copies of the parent's set of open file descriptors
 - ♦ As well as status flags and current file offset

Example 2a

- Two processes writing something to the standard output

```
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>

void char_at_a_time( const char * str ) {
    while( *str!= '\0' ) {
        putchar( *str++ );    // Write a char and increment the pointer
        fflush( stdout );     // Print out immediately (no buffering)
        usleep(50);
    }
}

int main() {
    if( fork() == 0 )
        char_at_a_time( "....." );
    else
        char_at_a_time( "||||||||||||" );
}
```

Example 2a

```
$ gcc forkme_sync1.cpp -o forkme
$ ./forkme
```

```
|.|.|.|.|.|.|.|.|.|.|.|.|.|.|
```

- Concurrency leads to unpredictable processes execution order
- The application might need to *synchronize* the execution of two or more processes
- The parent process might need to *wait for* a child process to finish
 - ♦ The parent process forks a child process to perform a computation, goes on in parallel, and then it reaches an execution point where it needs to use the output data of the child process
- Considering our example, assume this is the output we want:

```
.....| | | | | | | | | |
```


Example 2b: *Forking a process with synchronization*

- The results can be obtained by exploiting `wait(...)` functions

```
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>

void char_at_a_time( const char * str ) {
    while( *str!= '\0' ) {
        putchar( *str++ );    // Write a char and increment the pointer
        fflush( stdout );     // Print out immediately (no buffering)
        usleep(50);
    }
}

int main() {
    if( fork() == 0 )
        char_at_a_time( "....." );
    else {
        wait( NULL );
        char_at_a_time( "||||||||||||" );
    }
}
```

Synchronization using `wait()`

- The parent process block itself until a status change has occurred in one of the child processes
 - ♦ Child process terminated or stopped
 - ♦ Child process resumed by a *signal* (see later)
- The status is retrieved from an integer argument passed by pointer
 - ♦ `pid_t wait(int * status)`
- The `waitpid(pid_t, ...)` call allows the caller process to wait for a specific child process
- The `wait/waitpid` calls allow the system to release the resources associated with the child process
 - ♦ (e.g., opened files, allocated memory, etc...)

Zombie processes

- If a child terminates, without `wait()` performed, it remains in a “zombie” state
- The Linux kernel maintains a minimal set of information about the zombie process
 - (PID, termination status, resource usage information, ...)
 - Parent can later perform a wait to obtain information about children
- A zombie process consumes a slot in the *kernel process table*
 - If the table fills, it will not be possible to create further processes
- If a parent process terminates, then its “zombie” children (if any) are adopted by the *init* process
 - *init* automatically performs a `wait` to remove the zombies

Spawning executor processes

- Process forking basically “clones” the parent process image
 - ♦ Same code and same variables
- In a multi-process application we may need to spawn a process to execute a completely different task (program)
 - ♦ Load and run another executable
- The **exec ()** family of functions allows us to start a program within another program
- The **exec ()** family of functions replace the current process image with a new one coming from loading a new program

Spawning executor processes

- Function signatures

```
int execl(const char *path, const char *arg, ... );  
int execlp(const char *file, const char *arg, ... );  
int execle(const char *path, const char *arg, ... ,char * const envp[]);  
int execv(const char *path, char *const argv[]);  
int execvp(const char *file, char *const argv[]);
```

- All the functions take the executable path as first argument
- “**l**” functions accept variable amount of null-terminated char *
- “**v**” functions accept the executable path and an array of null-terminated char *
 - Both forward arguments to the executable (arg0 must be set to executable name)
- “**p**” functions access PATH environment variable to find the executable
- “**e**” functions accept also an array of null-terminated char * storing environment variables

Example 3

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>

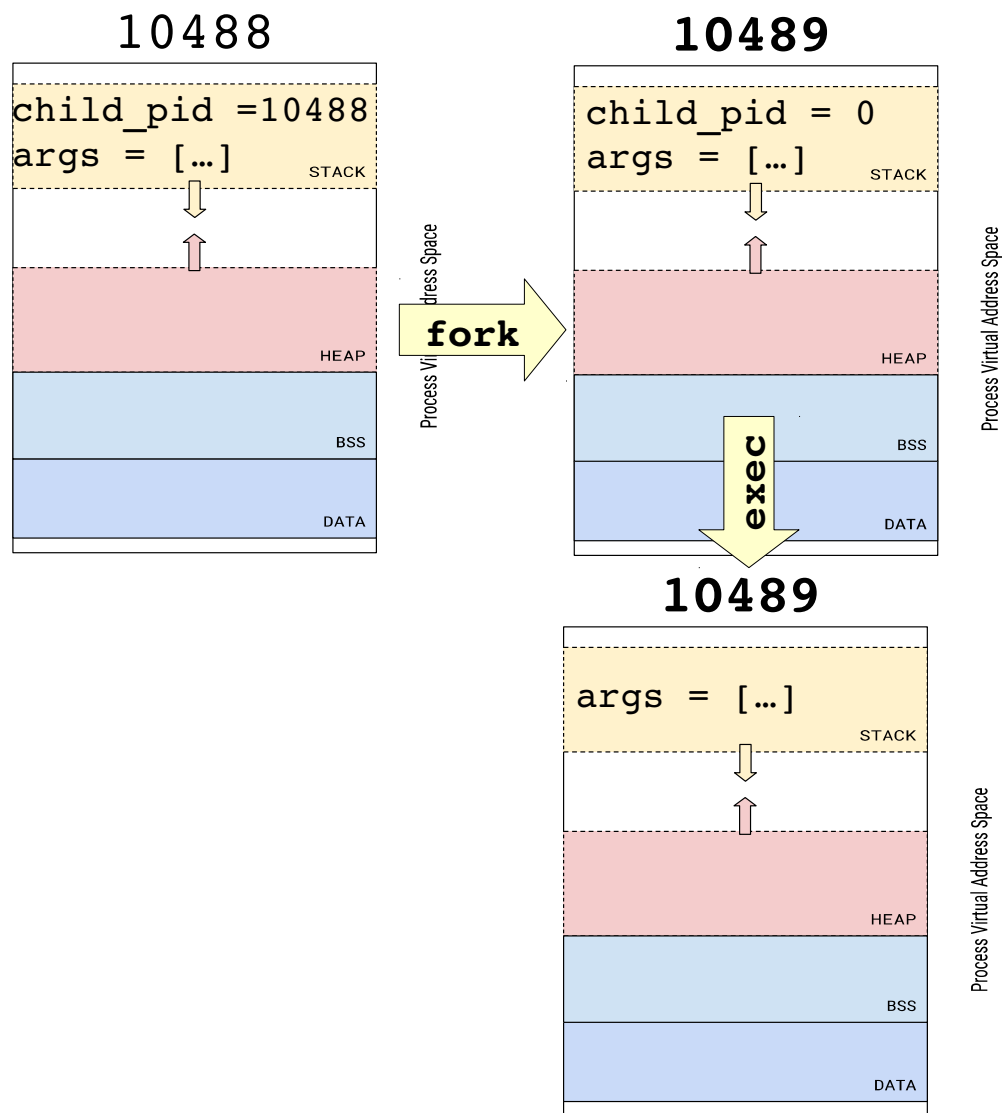
int spawn(const char * program, char ** arg_list) {
    pid_t child_pid = fork();
    if (child_pid != 0)
        return child_pid;          /* This is the parent process. */
    else {
        execvp (program, arg_list); /* Now execute PROGRAM */
        fprintf (stderr, "An error occurred in execvp\n");
        abort ();
    }
}

int main() {
    char * arg_list[] = { "ls", "-l", "/", NULL };
    spawn("ls", arg_list);
    printf ("Main program exiting...\n");
    return 0;
}
```

Spawning executor process

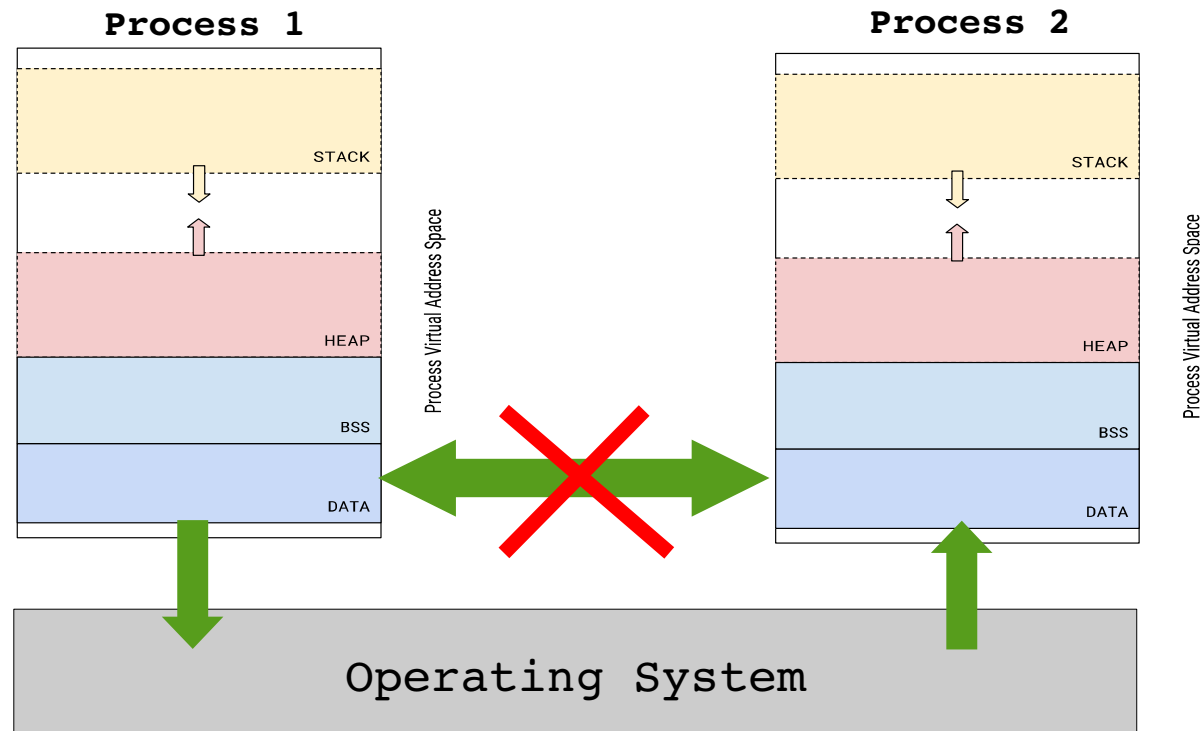
```
int main ()
{
    char * args[] = {
        "ls", "-l", NULL
    };
    pid_t child_pid;

    child_pid = fork();
    execvp("ls", arg);
    ...
}
```



Overview

- Each process has its own address space → How can we exchange information between different processes?
- Operating systems provide system calls on top of which communication mechanisms and API are built



Characteristics

- A single bit length “message”
- No data exchange
- Information content implicitly provided by the signal type
- Mechanism for asynchronous event notification

Examples

- Elapsed timer
- I/O operation completion
- Program exceptions
- User-defined events

Synchronization

- Asynchronous interaction between a sender and a receiver

Signal handling

- Signals may be thought as software equivalent of *hardware interrupts*
- Operating Systems manage a *signal vector table* for each process
 - ♦ Conversely, for hardware interrupts there is a single system-wide table
- OS typically defines several signals (name defined as integer macro)
- In Linux, the default action performed to handle a signal is to *terminate* the process
- It is possible to register a custom *signal handler* for each signal
 - ♦ Each entry of the signal vector table are
- Signals can be *ignored* at process-level (completely discarded)
- Signals can be *blocked* at process or thread-level
 - ♦ Enqueued and managed later, when the process/thread “unmask” the signal

A subset of common POSIX signals

POSIX signals	Portable number	Default action	Description
SIGABRT	6	Terminate	Process abort signal
SIGALRM	14		Alarm clock
SIGCHLD	N/A	Ignore	Child process terminated, stopped or continued
SIGINT	2	Terminate	Terminal interrupt
SIGKILL	9	Terminate	Kill the process
SIGPIPE	N/A	Terminate	Write on a pipe with no one to read it
SIGSEV	N/A	Terminate	Invalid memory reference
SIGUSR1	N/A	Terminate	User-defined signal 1
SIGUSR2	N/A	Terminate	User-defined signal 2
...

Example 4: User-defined signal handling

```
#include <signal.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <unistd.h>
sig_atomic_t sigusr1_count = 0;

void handler (int signal_number) {
    ++sigusr1_count;
}

int main() {
    struct sigaction sa;
    memset(&sa, 0, sizeof(sa));
    sa.sa_handler = &handler;
    sigaction (SIGUSR1, &sa, NULL);
    fprintf(stderr, "Running process... (PID=%d)\n", (int) getpid());
    /* Do some lengthy stuff here. */
    printf ("SIGUSR1 was raised %d times\n", sigusr1_count);
    return 0;
}
```

Example 4: *User-defined signal handling*

- Include `<signal.h>` header file
- Declare a data structure of type **sigaction**
- Clear the **sigaction** data structure and then set **sa_handler** field to point to the **handler()** function
- Register the signal handler for signal SIGUSR1 by calling the **sigaction()** function


```
$ gcc example4.cpp -o sig_example  
$ ./sig_example  
Running process... (PID=16151)
```

```
$ kill -SIGUSR1 16151
```

```
SIGUSR1 was raised 1 times
```

Unnamed pipes

- Based on the producer/consumer pattern
 - ♦ A producer write, a consumer read
- Data are written/read in a First-In First-Out (FIFO) fashion



- In Linux, the operating system guarantees that only one process per time can access the pipe
- Data written by the producer (sender) are stored into a buffer by the operating system until a consumer (receiver) read it

Example 5: Simple unnamed pipe based messaging (1/2)

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

/* Write COUNT copies of MESSAGE to STREAM, pausing for a second
   between each.  */
void writer(const char * message, int count, FILE * stream) {
    for(; count > 0; --count) {
        fprintf(stream, "%s\n", message);
        fflush(stream);
        sleep(1);
    }
}

void reader(FILE * stream) {
    char buffer[1024];
    /* Read until we hit the end of the stream.  fgets reads until
       either a newline or the end-of-file.  */
    while(!feof(stream) && !ferror(stream)
           && fgets(buffer, sizeof(buffer), stream) != NULL)
        fputs(buffer, stdout);
}
```

Example 5: Simple unnamed pipe based messaging (2/2)

```
int main () {
    FILE * stream;
    /* Create pipe place the two ends pipe file descriptors in fds */
    int fds[2];

    pipe(fds);
    pid_t pid = fork();
    if(pid == (pid_t) 0) {    /* Child process (consumer) */
        close(fds[1]);        /* Close the copy of the fds write end */
        stream = fdopen(fds[0], "r");
        reader(stream);
        close(fds[0]);
    }
    else {                    /* Parent process (producer) */
        close(fds[0]);        /* Close the copy of the fds read end */
        stream = fdopen(fds[1], "w");
        writer("Hello, world.", 3, stream);
        close(fds[1]);
    }
    return 0;
}
```

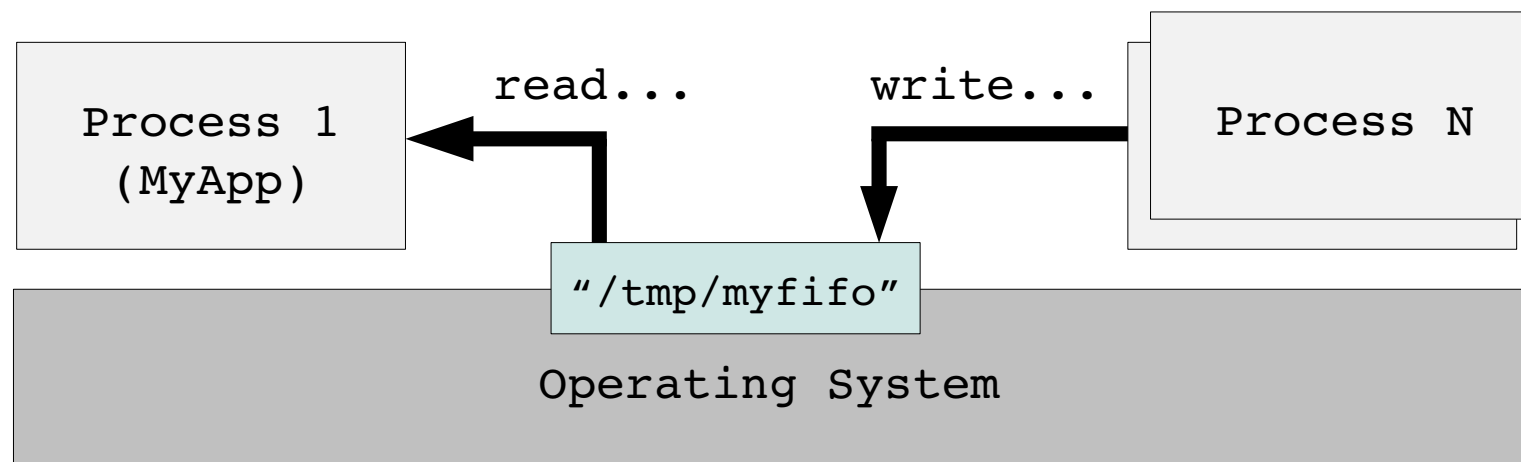

Example 5: Simple unnamed pipe based messaging

- Create a pipe with `pipe()` call and initialize the array of file descriptors “fds”
- Fork a child process that will behave as consumer
 - ♦ Close the write end of the pipe file descriptors array
 - ♦ Open the read end of the pipe file descriptors array
 - ♦ Call the `reader()` function to read data from the pipe
- Parent process acts as producer
 - ♦ Close the read end of the pipe file descriptors array
 - ♦ Open the write end of the pipe file descriptors array
 - ♦ Call the `writer()` function to write 3 times “*Hello, world.*”

```
Hello, world.  
Hello, world.  
Hello, world.
```

Named pipes (FIFO)

- Pipe-based mechanism accessible through file-system
- The pipe appears as a special FIFO file
- The pipe must be opened on both ends (reading and writing)
- OS passes data between processes without performing real I/O
- Suitable for *unrelated processes* communication



Example 6a: External interfacing through named pipe

▪ *fifo_writer.c*

```
int main () {
    struct datatype data;
    char * myfifo = "/tmp/myfifo";
    if (mkfifo(myfifo, S_IRUSR | S_IWUSR) != 0)
        perror("Cannot create fifo. Already existing?");

    int fd = open(myfifo, O_RDWR);
    if (fd == 0) {
        perror("Cannot open fifo");
        unlink(myfifo);
        exit(1);
    }
    int nb = write(fd, &data, sizeof(struct datatype));
    if (nb == 0)
        fprintf(stderr, "Write error\n");

    close(fd);
    unlink(myfifo);
    return 0;
}
```

Example 6a: *External interfacing through named pipe*

- *fifo_reader.c*

```
int main () {
    struct datatype data;
    char * myfifo = "/tmp/myfifo";

    int fd = open(myfifo, O_RDONLY);
    if (fd == 0) {
        perror("Cannot open fifo");
        unlink(myfifo);
        exit(1);
    }

    read(fd, &data, sizeof(struct datatype));
    ...

    close(fd);
    unlink(myfifo);
    return 0;
}
```

Example 6a: *External interfacing through named pipe*

- The writer
 - Creates the named pipe (`mkfifo`)
 - Open the named pipe as a normal file in read/write mode (`open`)
 - Write as many bytes as the size of the data structure
 - The reader must be in execution (otherwise data are sent to no nobody)
 - Close the file (`close`) and then release the named pipe (`unlink`)
- The reader
 - Open the named pipe as a normal file in read only mode (`open`)
 - The `read()` function blocks waiting for bytes coming from the writer process
 - Close the file (`close`) and then release the named pipe (`unlink`)

Example 6b: External interfacing through named pipe

- *message-reader.c*

- ♦ message-writer: the user sends char strings from the shell

```
int main () {
    char data = ' ';
    char * myfifo = "/tmp/myfifo";

    int fd = open(myfifo, O_RDWR);
    if (fd == 0) {
        perror("Cannot open fifo");
        unlink(myfifo);
        exit(1);
    }
    while (data != '#') {
        while (read(fd, &data, 1) && (data != '#'))
            fprintf(stderr, "%c", data);
    }
    close(fd);
    unlink(myfifo);
    return 0;
}
```

Example 6: *External interfacing through named pipe*

```
$ gcc example7.cpp -o ex_npipe
$ ./ex_npipe
Hello!
My name is
Joe
Communication closed
```

```
$ echo "Hello!" > /tmp/myfifo
$ echo "My name is" > /tmp/myfifo
$ echo "Joe" > /tmp/myfifo
$ echo "#" > /tmp/myfifo
```

- The (a priori known) named pipe location is opened as a regular file (`open`) to read and write
 - ♦ Write permission is required to flush data from pipe as they are read
- Blocking `read()` calls are performed to fetching data from the pipe
- The length of the text string not known a priori
 - ♦ ' #' is used as special END character
- Close (`close`) and release the pipe (`unlink`) when terminate

Pros

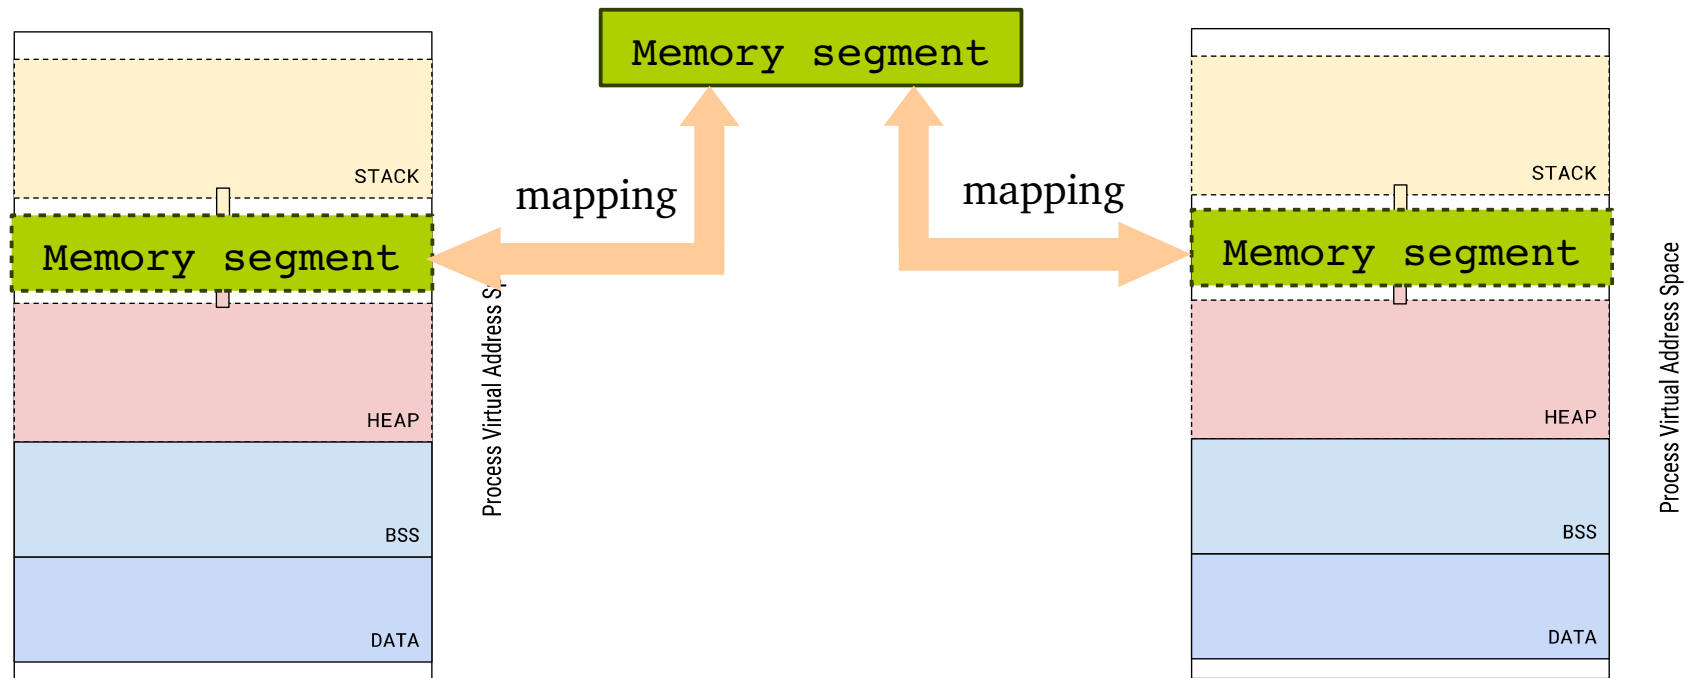
- Low overhead
- Simplicity
- Mutual access solved in kernel-space

Cons

- No broadcast
- Unidirectional
- No message boundaries, data are managed as a stream
- Poor scalability

Memory mapping

- Shared memory in Linux/UNIX operating systems is based on the concept of memory mapping
- A memory segment can be memory mapped in the address space of multiple processes



Memory mapping

- A POSIX standard has been defined to implement memory mapping application program interfaces
- **shm_open()** – opening/creation of a shared memory segment referenced by a name
 - A special file will appear in the file-system under “/dev/shm/” with the provided name
 - The special file represents a POSIX object and it is created for persistence
 - **ftruncate(...)** function resize to memory region to the correct size
- **mmap()** – mapping of the memory segment referenced by the file descriptor returned by **shm_open()**
- **munmap()** – unmapping of the memory segment
- **shm_unlink()** – removal of shared memory segment object if nobody is referencing it
- Link to POSIX real-time extension library to build (gcc ... **-lrt**)

Example 7: Simple shared memory mapping

♦ posix-shm-server.c (1/2)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main (int argc, char *argv[]) {
    const char * shm_name = "/AOS";
    const int SIZE = 4096;
    const char * message[] = {"This ", "is ", "about ", "shared ", "memory"};
    int i, shm_fd;
    void * ptr;
    shm_fd = shm_open(shm_name, O_CREAT | O_RDWR, 0666);
    if (shm_fd == -1) {
        printf("Shared memory segment failed\n");
        exit(-1);
    }
    ...
}
```

Example 7: Simple shared memory mapping

♦ posix-shm-server.c (2/2)

```
ftruncate(shm_fd, sizeof(message));
ptr = mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);
if (ptr == MAP_FAILED) {
    printf("Map failed\n");
    return -1;
}
/* Write into the memory segment */
for (i = 0; i < strlen(*message); ++i) {
    sprintf(ptr, "%s", message[i]);
    ptr += strlen(message[i]);
}
mummap(ptr, SIZE);
return 0;
}
```

- The server creates a shared memory referenced by “/AOS”
- The server writes some data (a string) into the memory segment
- The pointer ptr is incremented after each char string writing

Example 7: Simple shared memory mapping

- posix-shm-client.c (1/2)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main (int argc, char *argv[]) {
    const char * shm_name = "/AOS";
    const int SIZE = 4096;
    int i, shm_fd;
    void * ptr;

    shm_fd = shm_open(shm_name, O_RDONLY, 0666);
    if (shm_fd == -1) {
        printf("Shared memory segment failed\n");
        exit(-1);
    }
}
```

Example 7: Simple shared memory mapping

♦ posix-shm-client.c (2/2)

```
...    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);
    if (ptr == MAP_FAILED) {
        printf("Map failed\n");
        return -1;
    }
    printf("%s", (char *) ptr);

    if (shm_unlink(shm_name) == -1) {
        printf("Error removing %s\n", shm_name);
        exit(-1);
    }
    return 0;
}
```

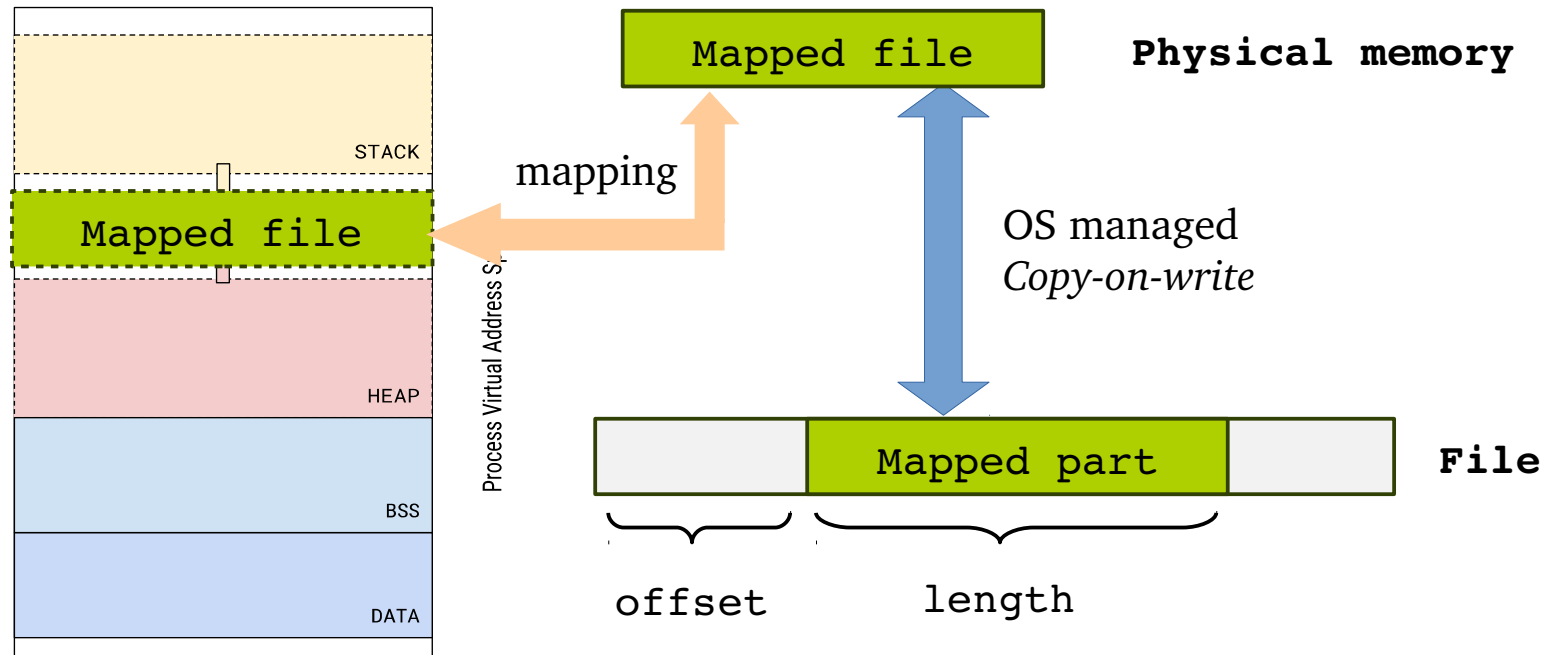
- The client opens the memory segment “AOS” in read-only mode
- The client maps the memory segment in read-only mode
- The client write the memory segment content to console

Example 7: *Simple shared memory mapping*

```
$ gcc posix-shm-server.c -o shm_server -lrt
$ gcc posix-shm-client.c -o shm_client -lrt
$ ./shm_server
$ ./shm_client
This is about shared memory
```

Memory mapping file

- Memory mapping allows us to logically insert part or all of a named *binary* file into a process address space



Example 8: Simple I/O mapping

- The file, passed at command-line (`argv[1]`), is opened and then memory mapped using the `mmap()` system call
- **`mmap()`** needs the address, the region size (file length), the permissions, the scope flags, file descriptor and offset

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <sys/mman.h>

int main (int argc, char *argv[]) {
    int * p;
    int fd = open(argv[1], O_RDWR);
    p = mmap(0, sizeof(int), PROT_READ | PROT_WRITE , MAP_SHARED, fd, 0);
    (*p)++;
    munmap(p, sizeof(int));
    close (fd);
    return 0;
}
```

Semaphores

- Concurrency in multi-tasking applications may introduce race conditions → we need to protect shared resource
- *Semaphores* are examples of structure aiming at solving such a problem in multi-process applications
- Semaphores are usually system objects managed by the OS kernel
- Semaphores can be thought as counters that we can manipulate by performing two actions: *wait* and *post*
- If counter value > 0 , *wait* decrements the counter and allows the task to enter the critical section
- If counter value = 0, *wait* blocks the tasks in a waiting list
- *post* increments the counter value
 - ♦ If the previous value was 0, a task is woken up from the waiting list

POSIX semaphores

- **sem_open()** – opening/creation of a *named* semaphore
 - ♦ Useful for synchronization among *unrelated processes*
- **sem_wait()** – Decrement the counter and lock if counter = 0
 - ♦ Initial counter value can be set to > 1
- **sem_post()** – Increment the count and unlock the critical section if counter > 0
- **sem_close()** – Close all the references to the named semaphore
- **sem_unlink()** – Destroy semaphore object
 - ♦ if all the references have been closed
- Link to POSIX real-time and threads extension library to build (gcc ... **-lrt -pthread**)

Example 9: Using semaphores with shared memory

▪ *posix-shm-sem-writer.c* (1/2)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <semaphore.h>

#define SHMOBJ_PATH  "/shm_AOS"
#define SEM_PATH     "/sem_AOS"

struct shared_data {
    char var1[10];
    int  var2;
};

int main(int argc, char *argv[]) {
    int shared_seg_size = (1 * sizeof(struct shared_data));
    ...
```

Example 9: Using semaphores with shared memory

▪ *posix-shm-sem-writer.c* (2/2)

```
int shmfd = shm_open(SHMOBJ_PATH, O_CREAT | O_RDWR, S_IRWXU | S_IRWXG);
ftruncate(shmfd, shared_seg_size);
struct shared_data * shared_msg = (struct shared_data *)
    mmap(NULL, shared_seg_size, PROT_READ | PROT_WRITE, MAP_SHARED,
        shmfd, 0);

sem_t * sem_id = sem_open(SEM_PATH, O_CREAT, S_IRUSR | S_IWUSR, 1);
struct shared_data out_msg = { "John", 23 };
sem_wait(sem_id);
/* Update shared data */
memcpy(shared_msg, &out_msg, sizeof(struct shared_data));
sem_post(sem_id);

shm_unlink(SHMOBJ_PATH);
sem_close(sem_id);
sem_unlink(SEM_PATH);
return 0;
}
```

Example 9: *Using semaphores with shared memory*

- The writer process
 - Maps a memory region
 - Creates a named semaphore and initialize it to 1 (`sem_open`)
 - Decrements the semaphore counter acquiring an exclusive access to the shared memory region (`sem_wait`)
 - Write into the memory region (`memcpy`)
 - Decrements the semaphore counter and releases the access to the memory region (`sem_post`)
 - Releases the shared memory region (`shm_unlink`)
 - Close and release the semaphore object (`sem_unlink`)

Example 9: Using semaphores with shared memory

▪ *posix-shm-sem-reader.c (1/2)*

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <semaphore.h>

#define SHMOBJ_PATH  "/shm_AOS"
#define SEM_PATH     "/sem_AOS"

struct shared_data {
    char var1[10];
    int  var2;
};

int main(int argc, char *argv[]) {
    int shared_seg_size = (1 * sizeof(struct shared_data));
    ...
```

Example 9: Using semaphores with shared memory

▪ *posix-shm-sem-reader.c (2/2)*

```
int shmfd = shm_open(SHMOBJ_PATH, O_RDONLY, 0666);

struct shared_data * shared_msg = (struct shared_data *)
    mmap(NULL, shared_seg_size, PROT_READ, MAP_SHARED, shmfd, 0);

sem_t * sem_id = sem_open(SEM_PATH, 0);
struct shared_data in_msg;
sem_wait(sem_id);
/* Update shared data */
memcpy(&in_msg, shared_msg, sizeof(struct shared_data));
sem_post(sem_id);
/* Process data... */

shm_unlink(SHMOBJ_PATH);
sem_close(sem_id);
sem_unlink(SEM_PATH);
return 0;
}
```


Example 9: *Using semaphores with shared memory*

- The reader process
 - Maps a memory region (read-only access)
 - Open the semaphore object, already initialized (`sem_open`)
 - Decrements the semaphore counter acquiring an exclusive access to the shared memory region (`sem_wait`)
 - Copy the data from the memory region to a local variable (`memcpy`)
 - Decrements the semaphore counter and releases the access to the memory region (`sem_post`)
 - Process the data
 - Releases the shared memory region (`shm_unlink`)
 - Close and release the semaphore object (`sem_unlink`)

Pros

- Can reduce memory usage

A big data structure can be mapped and shared to provide the same input set to multiple processes

- I/O mapping can be very efficient
 - ♦ Memory accesses instead of I/O read/write
 - ♦ Memory pages written back by OS only if the content has been modified
 - ♦ Seeking into the file performed with pointer arithmetic instead of “`lseek`”

Cons

- Linux map memory with a granularity of *memory page size*
 - ♦ Linux memory page size is typically 4 KB
 - ♦ Use memory mapping to map big files or share big data structures
- Can lead to *memory fragmentation*
 - ♦ Especially on 32-bits architectures
- Multiple small mappings can weight in terms of OS overhead